Efficient Multimodal Transformer with Dual-Level Feature Restoration for Robust Multimodal Sentiment Analysis

Licai Sun[1], Zheng Lian[2], Bin Liu[2], Jianhua Tao[3]

[1]UCAS & [2]CASIA & [3]Tsinghua University

How to fuse unaligned multimodal sequence?

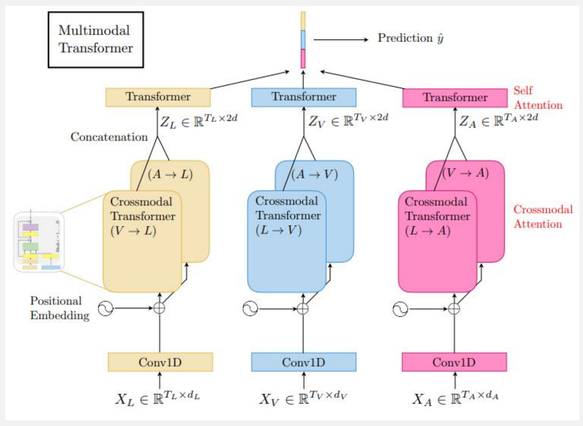

MulT (ACL 2019)
One-to-One Local-Local Fusion
Complexity: $O(M^2T^2)$

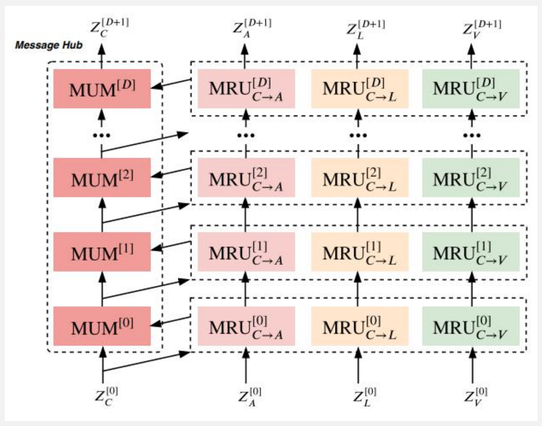

PMR (CVPR 2021)
One-to-All Local-Local Fusion
Complexity: $O(M^2T^2)$

How to handle random feature missing?

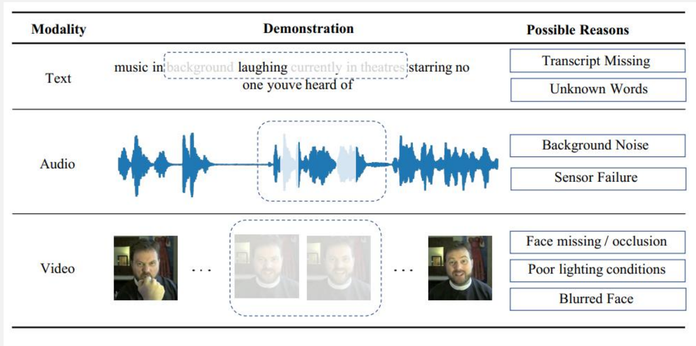

TFR-Net (MM 2021)
Low-Level Feature Reconstruction

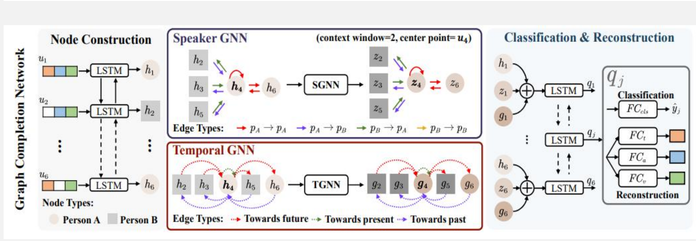

GCNet (TPAMI 2023)
Low-Level Feature Reconstruction

Our solution: EMT-DLFR

EMT

One-to-All Global-Local Fusion
Complexity: $O(MT^2)$

$$H_l^{[i+1]}, G_{l\to g}^{[i]} = \mathrm{MPU}_{l\leftrightarrow g}^{[i]}(H_l^{[i]}, G^{[i]})$$
$$H_a^{[i+1]}, G_{a\to g}^{[i]} = \mathrm{MPU}_{a\leftrightarrow g}^{[i]}(H_a^{[i]}, G^{[i]})$$
$$H_v^{[i+1]}, G_{v\to g}^{[i]} = \mathrm{MPU}_{v\leftrightarrow g}^{[i]}(H_v^{[i]}, G^{[i]})$$

MPU

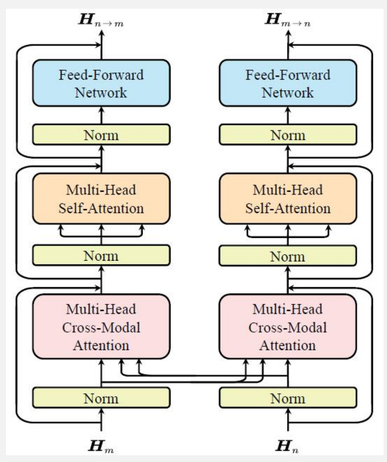

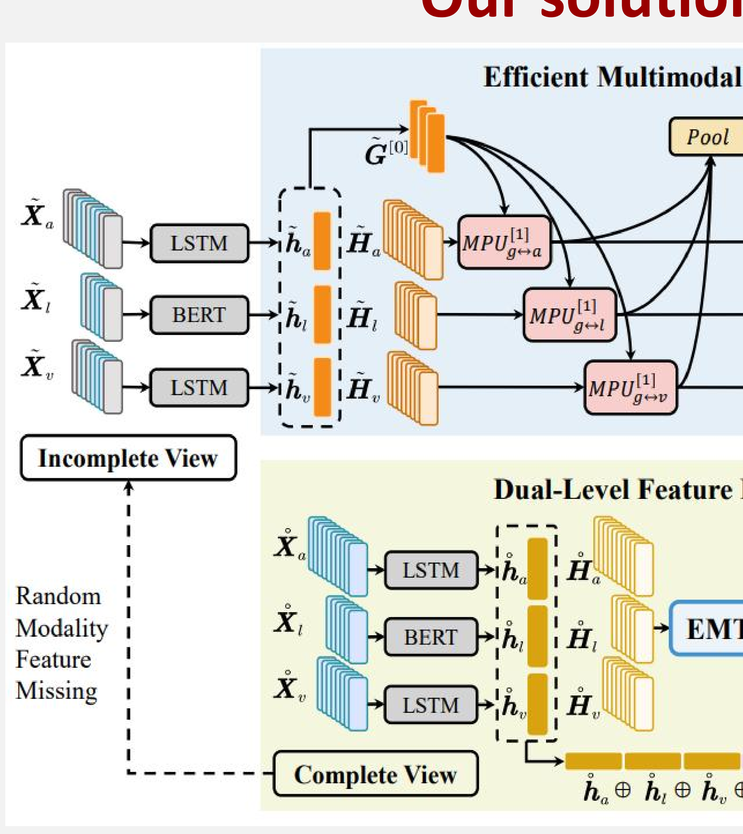

DLFR

Low-Level Feature Reconstruction (LLFR)
+
High-Level Feature Attraction (HLFR)

$$\mathcal{L}_{\text{recon}}^{l} = \text{smooth}_{\text{L1}}((\mathring{H}_l - r(\tilde{Z}_l)) \cdot (1 - g_l))$$
$$\mathcal{L}_{\text{recon}}^{a} = \text{smooth}_{\text{L1}}((\mathring{X}_a - r(\tilde{Z}_a)) \cdot (1 - g_a))$$
$$\mathcal{L}_{\text{recon}}^{v} = \text{smooth}_{\text{L1}}((\mathring{X}_v - r(\tilde{Z}_v)) \cdot (1 - g_v))$$
$$\mathcal{L}_{\text{recon}} = \sum_{m\in\{l,a,v\}} \mathcal{L}_{\text{recon}}^{m}$$

: Missing Feature
: Complete Feature
: Incomplete View
: Complete View
l : Text Modality
a : Audio Modality
v : Vision Modality
$\oplus$: Concatenation

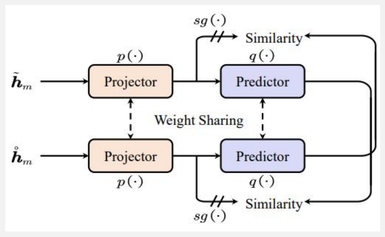

SimSiam for HLFR

Quantitative results

Complete setting

Models	CMU-MOSI MAE (↓)	Corr (↑)	Acc-7 (↑)	Acc-5 (↑)	Acc-2 (↑)	F1 (↑)
TFN†	0.901	0.698	34.9	-	-/80.8	-/80.7
LMF†	0.917	0.695	33.2	-	-/82.5	-/82.4
MulT†	0.861	0.711	-	-	81.5/84.1	80.6/83.9
MISA†	0.804	0.764	-	-	80.8/82.1	80.8/82.0
Self-MM†	0.712	0.795	45.8	-	82.5/84.8	82.7/84.9
MMIM†	0.700	0.800	46.7	-	84.1/86.1	84.0/86.0
AMML‡	0.723	0.792	46.3	-	-/84.9	-/84.8
TFR-Net◇	0.754	0.783	-	54.7	-/84.1	-/-
MulT	0.846	0.725	40.4	46.7	81.7/83.4	81.9/83.5
Self-MM	0.717	0.793	46.4	52.8	82.9/84.6	82.8/84.6
MMIM	0.712	0.790	46.9	53.0	**83.3/85.3**	**83.4/85.4**
TFR-Net	0.721	0.789	46.1	53.2	82.7/84.0	82.7/84.0
EMT	**0.705**	**0.798**	**47.4**	**54.1**	**83.3**/85.0	83.2/85.0

Incomplete setting

Models	CMU-MOSI MAE (↓)	Corr (↑)	Acc-7 (↑)	Acc-5 (↑)	Acc-2 (↑)	F1 (↑)
TFN◇	1.327	0.300	-	23.3	-/60.4	-/-
MulT◇	1.288	0.334	-	24.4	-/61.8	-/-
MISA◇	1.209	0.403	-	27.1	-/63.2	-/-
TFR-Net◇	1.155	0.467	-	30.4	-/69.0	-/-
TFN	1.316	0.308	22.3	23.7	61.0/60.9	59.7/59.7
LMF	1.310	0.299	21.5	22.7	59.7/59.3	56.4/56.1
MulT	1.263	0.348	23.1	24.6	63.1/63.2	60.7/61.0
MISA	1.202	0.405	25.7	27.4	63.9/63.7	59.0/58.8
Self-MM	1.162	0.444	27.8	30.3	66.9/67.5	65.4/66.2
MMIM	1.168	0.450	27.0	29.4	66.8/66.9	64.6/65.8
TFR-Net	1.156	0.452	27.7	30.5	67.6/67.8	65.7/66.1
EMT-DLFR	**1.106**	**0.486**	**32.5**	**35.6**	**69.6/70.3**	**69.6/70.3**

EMT achieves on-par/better performance **EMT-DLFR** achieves much better performance

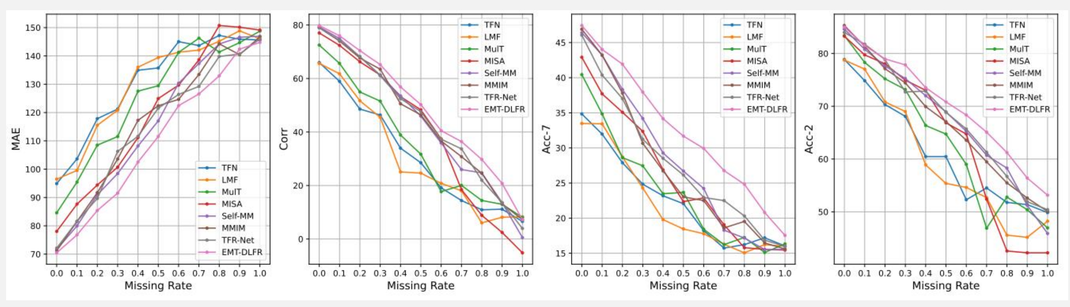

Ablation study

Fusion Strategy	MACs (G)	#Params (M)	Training Time (s)	GPU Memory (GB)
MulT	-	111.0	17.5	17.8
TFR-Net	-	124.3	24.8	16.9
OOLL	3.1	**110.5**	17.1	17.8
OALL	8.3	**110.5**	21.2	31.5
OAGL	**1.5**	**110.5**	**15.4**	**10.8**

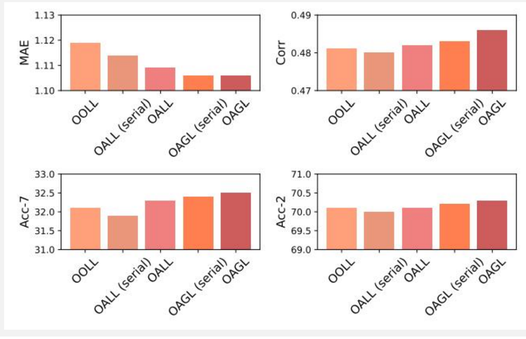

- **EMT** is *effective* and *efficient*!
- **HLFR** is more *effective* than LLFR and they are *complementary*!

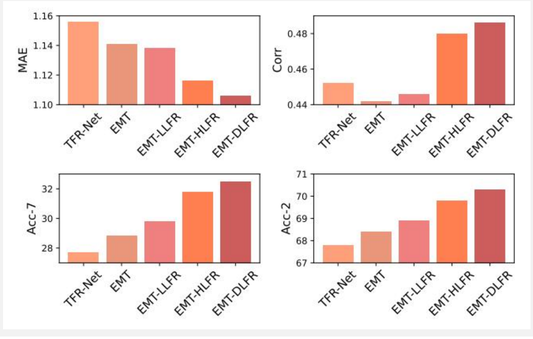

Qualitative results

- **EMT-DLFR** is *robust* to random feature missing!

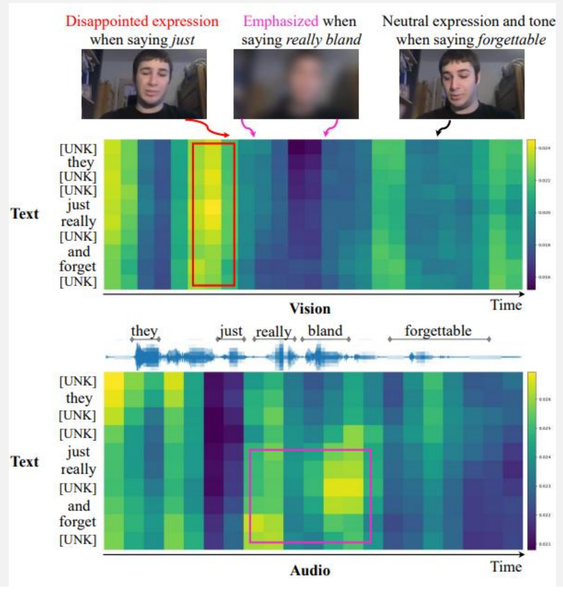

- Previous local-local fusion is *low-rank* and *redundant*!

More information can be found in our paper and code:

Paper:

Code: